


# **TensorFlow: A Framework for Scalable Machine Learning**



**You probably  
want to know...**

**What is TensorFlow?**

**Why did we create TensorFlow?**

**How does Tensorflow Work?**

**Example: Linear Regression**

**Example: Convolutional Neural Network**

**Distributed TensorFlow**






Fast, flexible, and scalable  
open-source machine learning  
library

One system for research and  
production

Runs on CPU, GPU, TPU, and  
Mobile

Apache 2.0 license

# TensorFlow Handles Complexity



**Modeling complexity**



**Distributed  
System**



**Heterogenous  
System**




## **Under the Hood**

A multidimensional array.



**Tensor**Flow



A graph of operations.

# The TensorFlow Graph

Computation is defined as a graph

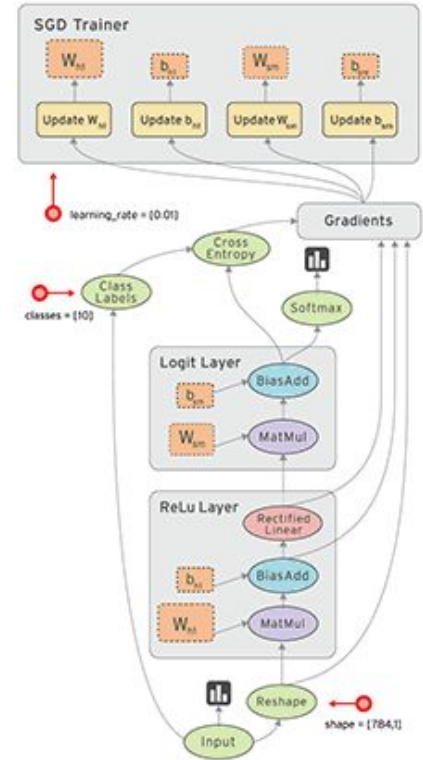
Graph is defined in high-level language (Python)

Graph is compiled and optimized

Graph is executed (in parts or fully) on available low level devices (CPU, GPU, TPU)

Nodes represent computations and state

Data (tensors) flow along edges



## Build a graph; then run it.


...

```
c = tf.add(a, b)
```

...


```
session = tf.Session()
```

```
value_of_c = session.run(c, {a=1, b=2})
```






# Any Computation is a TensorFlow Graph




# Any Computation is a TensorFlow Graph

*with state*



# Automatic Differentiation


**Automatically add ops which  
compute gradients for variables**



# Any Computation is a TensorFlow Graph


Simple gradient descent:

**with state**



# Any Computation is a TensorFlow Graph

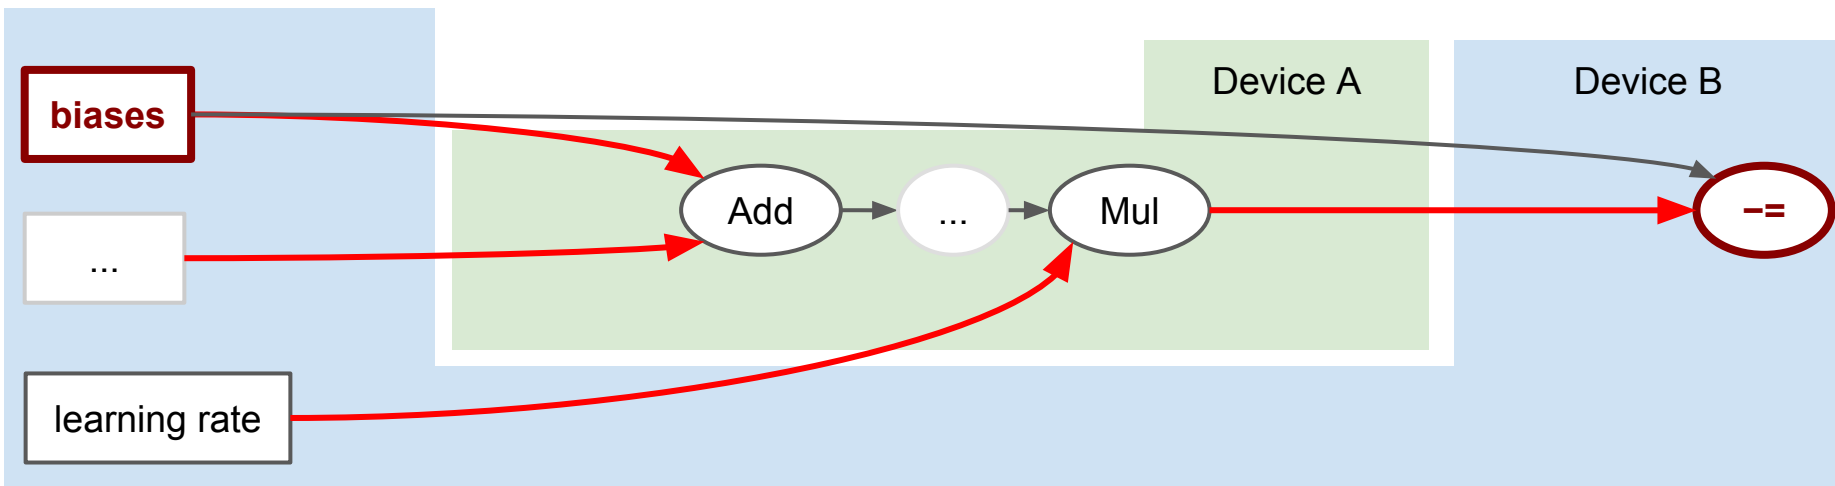
**distributed**



Devices: Processes, Machines, CPUs, GPUs, TPUs, etc

# Send and Receive Nodes

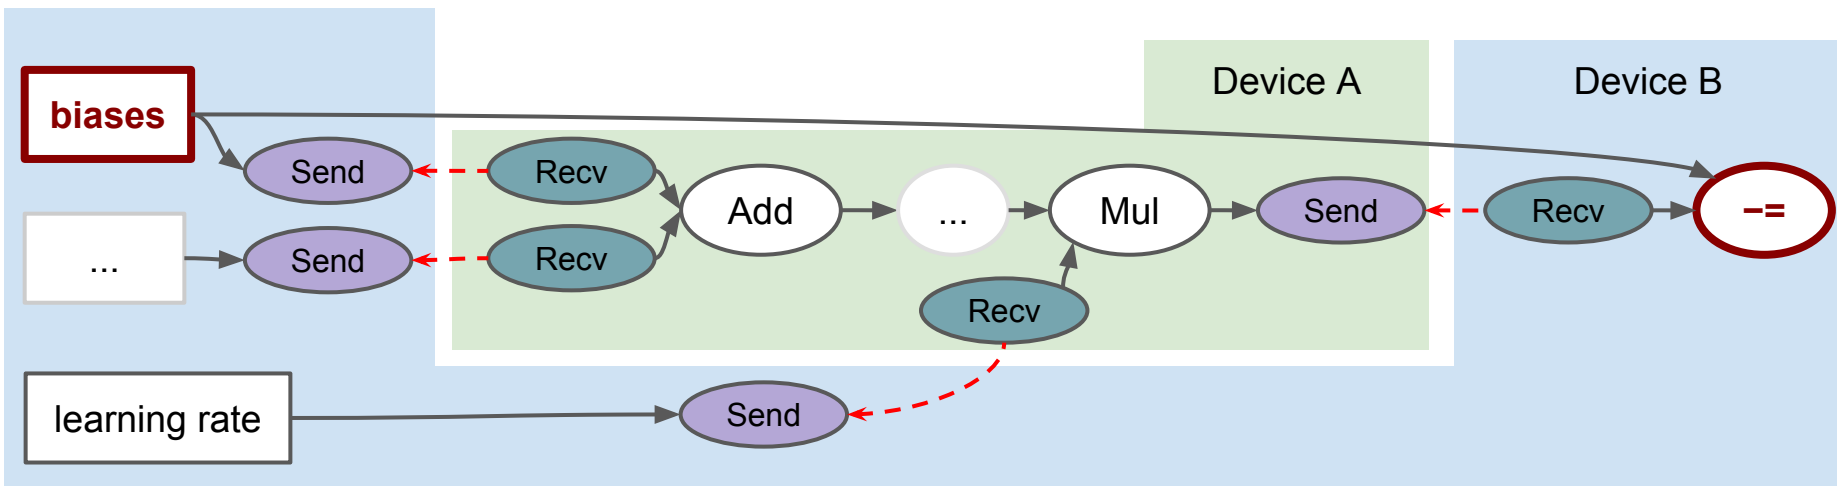
**distributed**



Devices: Processes, Machines, CPUs, GPUs, TPUs, etc

# Send and Receive Nodes

**distributed**



Devices: Processes, Machines, CPUs, GPUs, TPUs, etc



# Linear Regression




# Linear Regression

Diagram illustrating the linear regression equation  $y = Wx + b$ .

The equation components are labeled with arrows:

- result** points to  $y$ .
- input** points to  $x$ .
- parameters** points to  $W$  and  $b$ .



# What are we trying to do?

**Mystery equation:**  $y = 0.1 * x + 0.3 + \text{noise}$

**Model:**  $y = W * x + b$

**Objective:** Given enough  $(x, y)$  value samples, figure out the value of  $W$  and  $b$ .

# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf

x = tf.placeholder(shape=[None],
                   dtype=tf.float32, name="x")
```

# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf

x = tf.placeholder(shape=[None],
                   dtype=tf.float32, name="x")

W = tf.get_variable(shape=[], name="W")
```

# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf

x = tf.placeholder(shape=[None],
                   dtype=tf.float32, name="x")

W = tf.get_variable(shape=[], name="W")

b = tf.get_variable(shape=[], name="b")
```


# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf

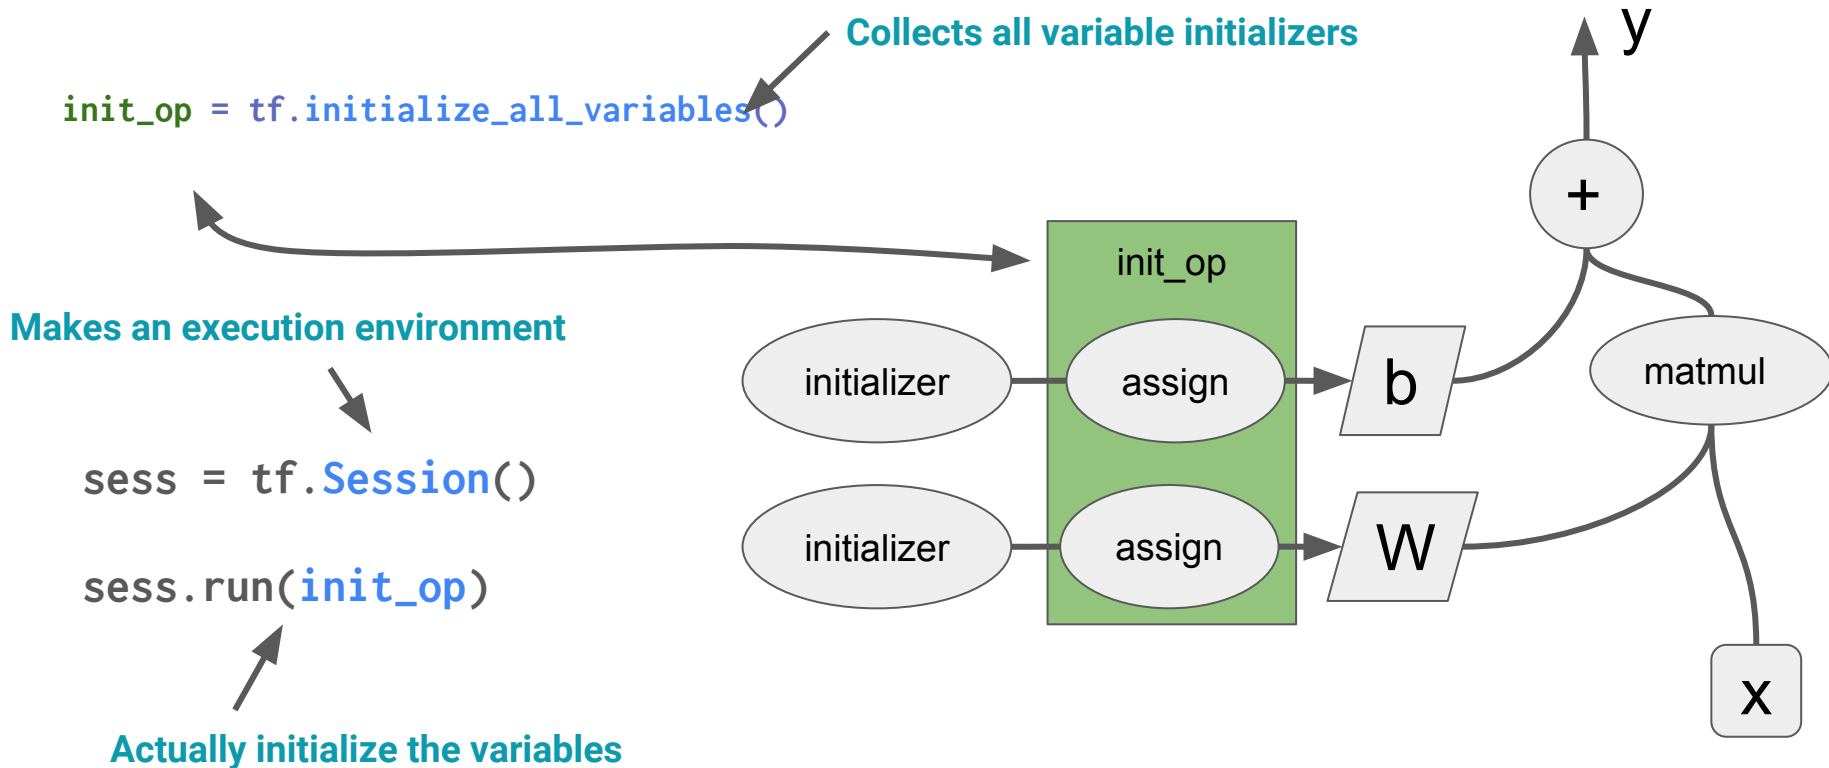
x = tf.placeholder(shape=[None],
                   dtype=tf.float32, name="x")

W = tf.get_variable(shape=[], name="W")
b = tf.get_variable(shape=[], name="b")

y = W * x + b
```



# Variables Must be Initialized






# Running the Computation

```
x_in = [3]
```

```
sess.run(y, feed_dict={x: x_in})
```

Only what's used to compute a fetch will be evaluated  
All Tensors can be fed, but all placeholders must be fed



# Putting it all together

```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                   dtype=tf.float32,
                   name='x')

W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b
```

```
with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())
    print(sess.run(y, feed_dict={x: x_in}))
```

} Build the graph

} Prepare execution environment

} Initialize variables

} Run the computation (usually often)


# Define a Loss

Given  $y$ ,  $y_{train}$  compute a loss, for instance:

$$L = (y - y_{train})^2$$

```
# create an operation that calculates loss.  
loss = tf.reduce_mean(tf.square(y - y_train))
```

## Minimize loss: optimizers



`tf.train.AdadeltaOptimizer`  
`tf.train.AdagradOptimizer`  
`tf.train.AdagradDAOptimizer`  
`tf.train.AdamOptimizer`  
...

# Train

Feed  $(x, y_{\text{label}})$  pairs and adjust  $W$  and  $b$  to decrease the loss.

$$W \leftarrow W - \eta (dL/dW)$$


$$b \leftarrow b - \eta (dL/db)$$

```
# Create an optimizer
```

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

```
# Create an operation that minimizes loss.
```

```
train = optimizer.minimize(loss)
```



TensorFlow computes  
gradients automatically



Learning rate

# Putting it all together

```
loss = tf.reduce_mean(tf.square(y - y_train))
```

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

```
train = optimizer.minimize(loss)
```

```
with tf.Session() as sess:
```

```
    sess.run(tf.initialize_all_variables())
```

```
    for i in range(1000):
```

```
        sess.run(train, feed_dict={x: x_in, y_label: y_in})
```

} Define a loss

} Create an optimizer

} Op to minimize the loss

} Initialize variables

} Iteratively run the training op

# Putting it all together

```
import tensorflow as tf
x = tf.placeholder(shape=[None], dtype=tf.float32, name='x')
W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b

loss = tf.reduce_mean(tf.square(y - y_train))

optimizer = tf.train.GradientDescentOptimizer(0.5)

train = optimizer.minimize(loss)

with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())

    for i in range(1000):
        sess.run(train, feed_dict={x: x_in, y_label: y_in})
```

Build the graph

Define a loss

Create an optimizer

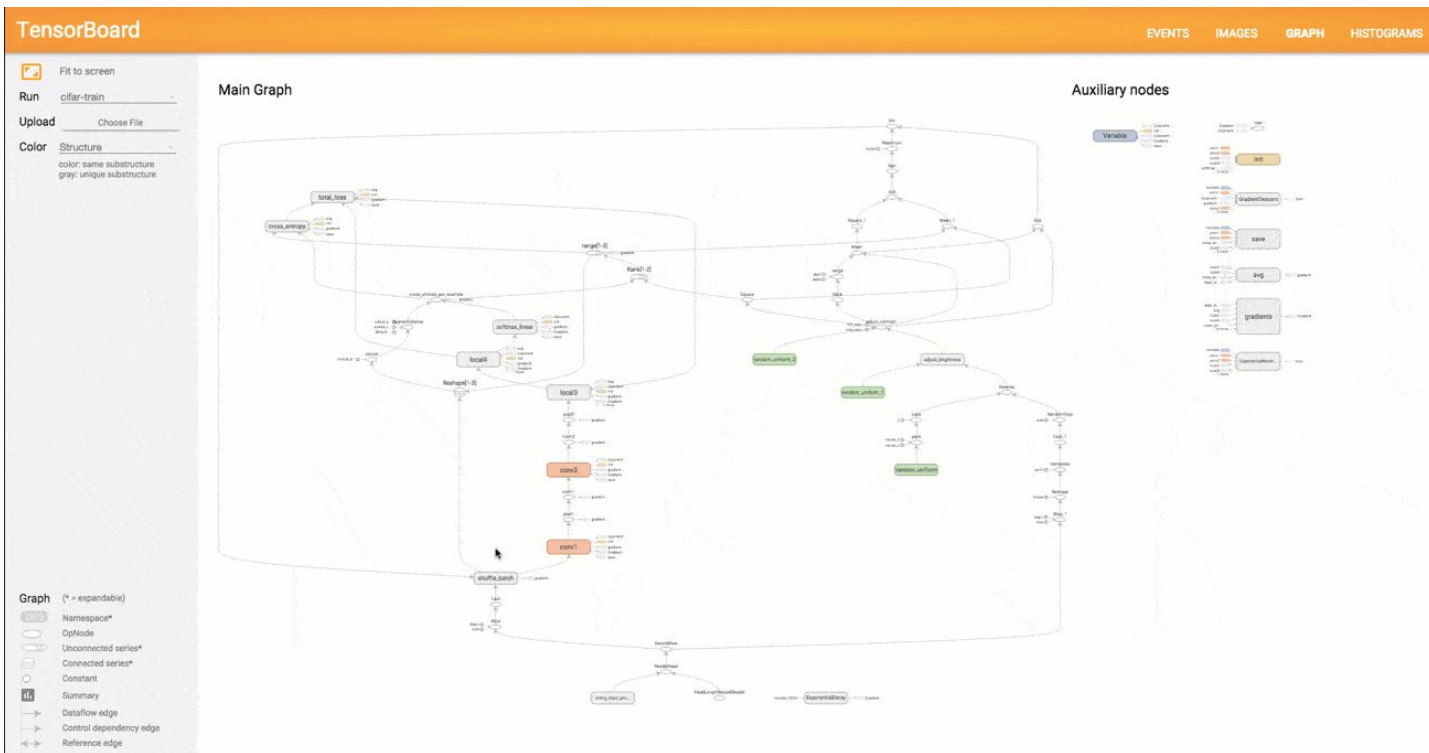
Op to minimize the  
loss

Prepare environment

Initialize variables

Iteratively run the  
training op

# TensorBoard








# **Deep Convolutional Neural Network**

IS THIS A  
**CAT or DOG?**



**CAT DOG**



# Remember linear regression?

```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                   dtype=tf.float32,
                   name='x')

W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b

loss = tf.reduce_mean(tf.square(y - y_label))

optimizer = tf.train.GradientDescentOptimizer(0.5)


train = optimizer.minimize(loss)

...
```


} Build the graph

# Convolutional DNN

```
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.fully_connected(x, activation_fn=tf.nn.relu)
x = tf.contrib.layers.dropout(x, 0.5)
logits = tf.config.layers.linear(x)
```




# Defining Complex Networks





# **Distributed TensorFlow**

# Data Parallelism (Between-Graph Replication)



# Describe a cluster: ClusterSpec

```
tf.train.ClusterSpec({  
    "worker": [  
        "worker0.example.com:2222",  
        "worker1.example.com:2222",  
        "worker2.example.com:2222"  
    ],  
    "ps": [  
        "ps0.example.com:2222",  
        "ps1.example.com:2222"  
    ]  
})
```



# Share the graph across devices

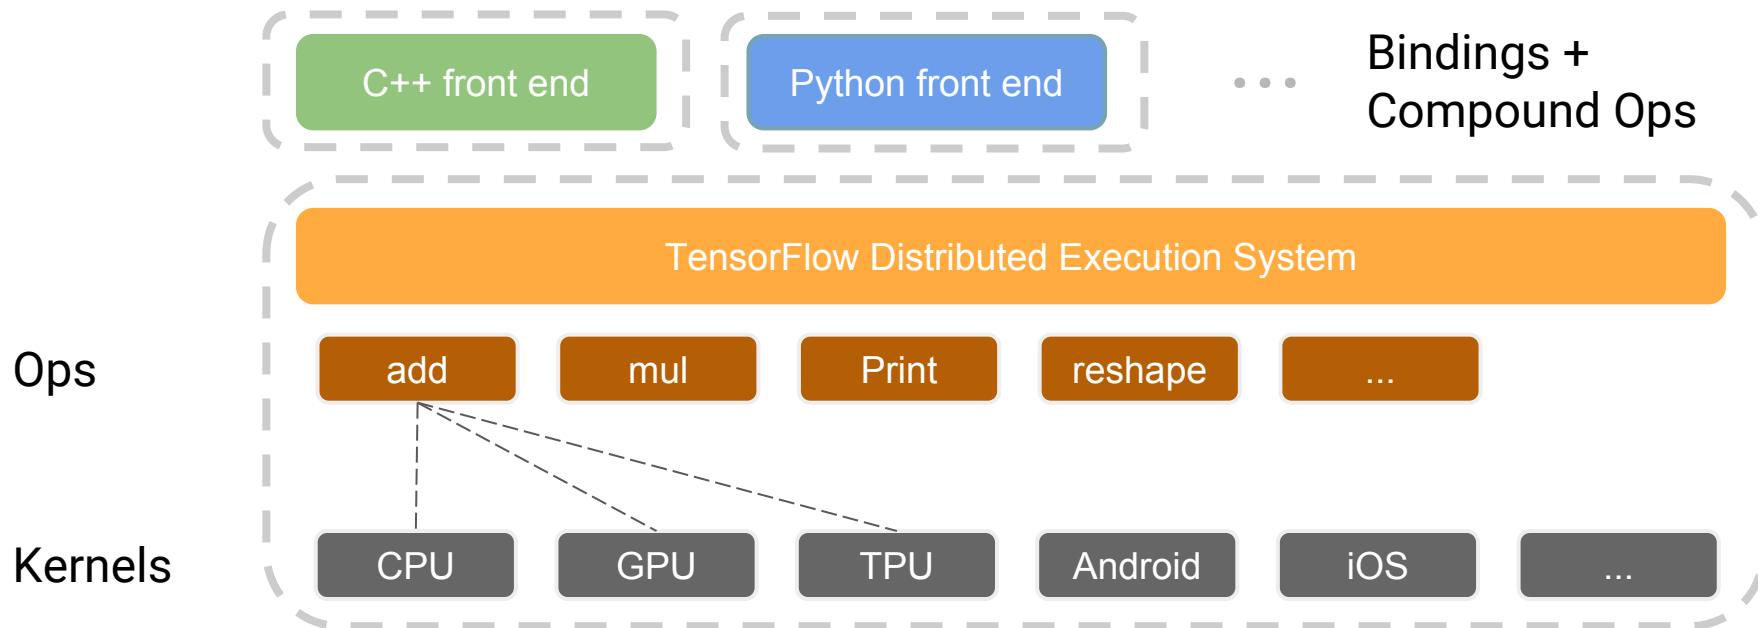
```
with tf.device("/job:ps/task:0"):
    weights_1 = tf.Variable(...)
    biases_1 = tf.Variable(...)
```

```
with tf.device("/job:ps/task:1"):
    weights_2 = tf.Variable(...)
    biases_2 = tf.Variable(...)
```

```
with tf.device("/job:worker/task:7"):
    input, labels = ...
    layer_1 = tf.nn.relu(tf.matmul(input, weights_1) + biases_1)
    logits = tf.nn.relu(tf.matmul(layer_1, weights_2) + biases_2)
    train_op = ...
```

```
with tf.Session("grpc://worker7.example.com:2222") as sess:
    for _ in range(10000):
        sess.run(train_op)
```

# Architecture



# Tutorials

**Tutorials on [tensorflow.org](https://www.tensorflow.org)**

**Image recognition:** [https://www.tensorflow.org/tutorials/image\\_recognition](https://www.tensorflow.org/tutorials/image_recognition)

**Word embeddings:** <https://www.tensorflow.org/tutorials/word2vec>

**Language Modeling:** <https://www.tensorflow.org/tutorials/recurrent>

**Translation:** <https://www.tensorflow.org/tutorials/seq2seq>

**Udacity Course:** <https://www.udacity.com/course/deep-learning--ud730>



**Q & A**



Google Developer Day